

```
! /usr/lib/wsl/lib/nvidia-smi
```

```

Sun Jul 13 16:28:58 2025
+-----+
+-----+
| NVIDIA-SMI 560.35.02                  Driver Version: 560.94          CUDA Version: 12.6
+-----+
+-----+
| GPU  Name                               Persistence-M | Bus-Id        Disp.A | Volatile Uncorr
r. ECC |
| Fan  Temp   Perf           Pwr:Usage/Cap |      Memory-Usage | GPU-Util  Compu
te M. |
|      |            |                               |                      |              M
IG M. |
+=====+
=====|
|   0  NVIDIA GeForce RTX 4090           On   |   00000000:01:00.0  On   |
Off |
| 30%   37C    P5              38W / 450W |   1897MiB / 24564MiB |      68%    De
fault |
|      |            |                               |                      |
N/A |
+-----+
+-----+

+-----+
+-----+
| Processes:
|
| GPU  GI    CI          PID    Type    Process name                      GPU M
emory |
|      ID    ID              |                       |              Usage
|
|=====+
=====|
|   0   N/A   N/A         31      G      /Xwayland                          N/A
|
+-----+
+-----+

```

```
# Imports
import os
import tensorflow as tf
from tensorflow import keras
from keras.callbacks import EarlyStopping
from keras import layers, models, regularizers
from keras.models import load_model
from keras.preprocessing import image
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from PIL import Image
```

2025-07-14 18:24:45.746603: I tensorflow/core/util/port.cc:113] oneDNN custom operations are on. You may see slightly different numerical results due to floating-point round-off errors from different computation orders. To turn them off, set the environment variable `TF_ENABLE_ONEDNN_OPTS=0`.

2025-07-14 18:24:45.897973: I tensorflow/core/platform/cpu_feature_guard.cc:210] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.

To enable the following instructions: SSE4.1 SSE4.2 AVX AVX2 AVX_VNNI FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.

Datenaufbereitung!!!!here!!!!

```
# 1. Datenaufbereitung
data_dir = "/home/cxy_otto/xuanyou/ABGA2/HA1/Data"
batch_size = 32
img_height = 224
```

```

img_width = 224
seed = 42

class_dirs = sorted([d for d in os.listdir(data_dir) if os.path.isdir(os.path.join(

class_stats = {} # Dictionary to store the number of images per class
image_sizes = set() # Set to store unique image sizes (to check for consistency)

for class_name in class_dirs:
    class_path = os.path.join(data_dir, class_name)
    images = [f for f in os.listdir(class_path) if f.lower().endswith((".jpg", ".pn
    class_stats[class_name] = len(images) # Store the number of images for this

    # first 3 images for size checking
    for img_file in images[:3]:
        with Image.open(os.path.join(class_path, img_file)) as img:
            image_sizes.add(img.size)

print("Klassen und Bildanzahl:")
for k, v in class_stats.items():
    # Print the number of images per class
    print(f" {k}: {v} Bilder")

print("\n Beispielhafte Bildgrößen:", set(image_sizes))

train_ds = keras.preprocessing.image_dataset_from_directory(
    data_dir,
    validation_split=0.2,
    subset="training",
    seed=seed,
    image_size=(img_height, img_width),
    batch_size=batch_size
)

val_ds = keras.preprocessing.image_dataset_from_directory(
    data_dir,
    validation_split=0.2,
    subset="validation",
    seed=seed,
    image_size=(img_height, img_width),
    batch_size=batch_size
)

class_names = train_ds.class_names
print(f"Name of class: {class_names}")
num_classes = len(class_names)
print(f"Number of classes: {num_classes}")

plt.figure(figsize=(15, 5))

for idx, class_name in enumerate(class_names):
    class_path = os.path.join(data_dir, class_name)
    image_files = sorted([f for f in os.listdir(class_path) if f.endswith((".jpg",
    if not image_files: # Skip if no image found in this class
        continue
    img = Image.open(os.path.join(class_path, image_files[0])) # Load the first
    plt.subplot(2, 5, idx + 1)
    plt.imshow(img)
    plt.title(class_name)
    plt.axis("off")

plt.suptitle("Beispielbilder pro Klasse")
plt.tight_layout()
plt.show()

# Prefetch + Rescaling
AUTOTUNE = tf.data.AUTOTUNE
train_ds = train_ds.cache().shuffle(1000).prefetch(buffer_size=AUTOTUNE)
val_ds = val_ds.cache().prefetch(buffer_size=AUTOTUNE)

```

Klassen und Bildanzahl:
Abstandshalter: 156 Bilder
Auslassventil: 156 Bilder
Blechlineal: 156 Bilder
Filterkartusche: 156 Bilder
Gewindestange: 156 Bilder
Hohlschraube: 156 Bilder
Hutmutter: 156 Bilder
Hydraulikstutzen: 156 Bilder
Nutenstein: 156 Bilder
Schraubenfeder: 156 Bilder

Beispielhafte Bildgrößen: {(2976, 2976), (3456, 3456), (1250, 1250), (3024, 3024)}

Found 1560 files belonging to 10 classes.

Using 1248 files for training.

Found 1560 files belonging to 10 classes.

Using 312 files for validation.

Name of class: ['Abstandshalter', 'Auslassventil', 'Blechlineal', 'Filterkartusche', 'Gewindestange', 'Hohlschraube', 'Hutmutter', 'Hydraulikstutzen', 'Nutenstein', 'Schraubenfeder']

Number of classes: 10



1. **Basismodell:** Das ursprüngliche Modell ($3 \times [\text{Conv}+\text{Pool}] + \text{Flatten} + \text{Dense}$) erzielte nach 40 Epochen eine Validierungsgenauigkeit (**val_acc**) von etwa **60,3 %** bei einem Verlust (**val_loss**) von **1,57**. Dieses Ergebnis dient als Referenzpunkt für die weitere Optimierung.

```
In [4]: # 2. Basismodell (einfach und kompakt)
model = models.Sequential([
    layers.Rescaling(1./255, input_shape=(img_height, img_width, 3)),
    layers.Conv2D(32, (3, 3), activation='relu'),
    layers.MaxPooling2D(2, 2),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D(2, 2),
    layers.Conv2D(128, (3, 3), activation='relu'),
    layers.MaxPooling2D(2, 2),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dropout(0.5),
    layers.Dense(num_classes, activation='softmax')
])
```

/home/cxy_otto/miniconda/envs/tf2-stable/lib/python3.11/site-packages/keras/src/layers/preprocessing/tf_data_layer.py:19: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.
super().__init__(**kwargs)

```
In [5]: # 3. Compile

model.compile(
    optimizer=keras.optimizers.Adam(learning_rate=1e-4),
    loss="sparse_categorical_crossentropy",
    metrics=["accuracy"]
)

# 4. Training (nur Basismodell)
```

```

history = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=40,
)

# 5. Evaluation: Training-Verläufe anzeigen
plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label='val_accuracy')
plt.legend()
plt.show()

loss, acc = model.evaluate(val_ds)
print(f"Validation accuracy: {acc:.4f} - Validation loss: {loss:.4f}")

```


Epoch 1/40

WARNING: All log messages before absl::InitializeLog() is called are written to STDERR


I0000 00:00:1752416949.605309 612510 service.cc:145] XLA service 0x748650007ab0 initialized for platform CUDA (this does not guarantee that XLA will be used). Devices: I0000 00:00:1752416949.605566 612510 service.cc:153] StreamExecutor device (0): NVIDIA GeForce RTX 4090, Compute Capability 8.9


2025-07-13 16:29:09.622103: I tensorflow/compiler/mlir/tensorflow/utils/dump_mlir_util.cc:268] disabling MLIR crash reproducer, set env var `MLIR_CRASH_REPRODUCER_DIRECTORY` to enable.


2025-07-13 16:29:09.722518: I external/local_xla/xla/stream_executor/cuda/cuda_dnn.cc:465] Loaded cuDNN version 8907


17/39  0s 10ms/step - accuracy: 0.1016 - loss: 2.3660


I0000 00:00:1752416952.464166 612510 device_compiler.h:188] Compiled cluster using XLA! This line is logged at most once for the lifetime of the process.


39/39  7s 44ms/step - accuracy: 0.1031 - loss: 2.3478 - val_accuracy: 0.1282 - val_loss: 2.2758
Epoch 2/40


39/39  0s 11ms/step - accuracy: 0.1770 - loss: 2.2342 - val_accuracy: 0.2340 - val_loss: 2.1663
Epoch 3/40


39/39  0s 11ms/step - accuracy: 0.2336 - loss: 2.0883 - val_accuracy: 0.2500 - val_loss: 2.0176
Epoch 4/40


39/39  0s 11ms/step - accuracy: 0.2848 - loss: 1.9341 - val_accuracy: 0.2917 - val_loss: 1.9069
Epoch 5/40


39/39  0s 11ms/step - accuracy: 0.3416 - loss: 1.7943 - val_accuracy: 0.3429 - val_loss: 1.8514
Epoch 6/40


39/39  0s 11ms/step - accuracy: 0.3905 - loss: 1.7073 - val_accuracy: 0.3654 - val_loss: 1.8147
Epoch 7/40


39/39  0s 11ms/step - accuracy: 0.3830 - loss: 1.6646 - val_accuracy: 0.3878 - val_loss: 1.7143
Epoch 8/40


39/39  0s 11ms/step - accuracy: 0.4483 - loss: 1.4980 - val_accuracy: 0.3878 - val_loss: 1.7253
Epoch 9/40


39/39  0s 11ms/step - accuracy: 0.4749 - loss: 1.4834 - val_accuracy: 0.4135 - val_loss: 1.6352
Epoch 10/40


39/39  0s 11ms/step - accuracy: 0.5282 - loss: 1.3776 - val_accuracy: 0.4231 - val_loss: 1.5895
Epoch 11/40


39/39  0s 11ms/step - accuracy: 0.5484 - loss: 1.2442 - val_accuracy: 0.5000 - val_loss: 1.5499
Epoch 12/40


39/39  0s 11ms/step - accuracy: 0.5893 - loss: 1.1854 - val_accuracy: 0.5160 - val_loss: 1.5310
Epoch 13/40


39/39  0s 11ms/step - accuracy: 0.6078 - loss: 1.1142 - val_accuracy: 0.4904 - val_loss: 1.5447
Epoch 14/40


39/39  0s 11ms/step - accuracy: 0.6386 - loss: 1.0474 - val_accuracy: 0.5160 - val_loss: 1.5383
Epoch 15/40


39/39  0s 11ms/step - accuracy: 0.6595 - loss: 1.0463 - val_accuracy: 0.5353 - val_loss: 1.4857
Epoch 16/40


39/39  0s 11ms/step - accuracy: 0.6727 - loss: 0.9459 - val_accuracy: 0.5353 - val_loss: 1.4529
Epoch 17/40


39/39  0s 11ms/step - accuracy: 0.7119 - loss: 0.8256 - val_accuracy: 0.5224 - val_loss: 1.5046
Epoch 18/40


39/39  0s 11ms/step - accuracy: 0.6946 - loss: 0.8898 - val_accuracy: 0.5673 - val_loss: 1.4077
Epoch 19/40


39/39  0s 11ms/step - accuracy: 0.7079 - loss: 0.8228 - val_accuracy: 0.5545 - val_loss: 1.4305
Epoch 20/40


39/39  0s 11ms/step - accuracy: 0.7449 - loss: 0.7511 - val_accuracy: 0.5705 - val_loss: 1.3937
Epoch 21/40

39/39  0s 11ms/step - accuracy: 0.7606 - loss: 0.7306 - val_accuracy: 0.5641 - val_loss: 1.4078
Epoch 22/40

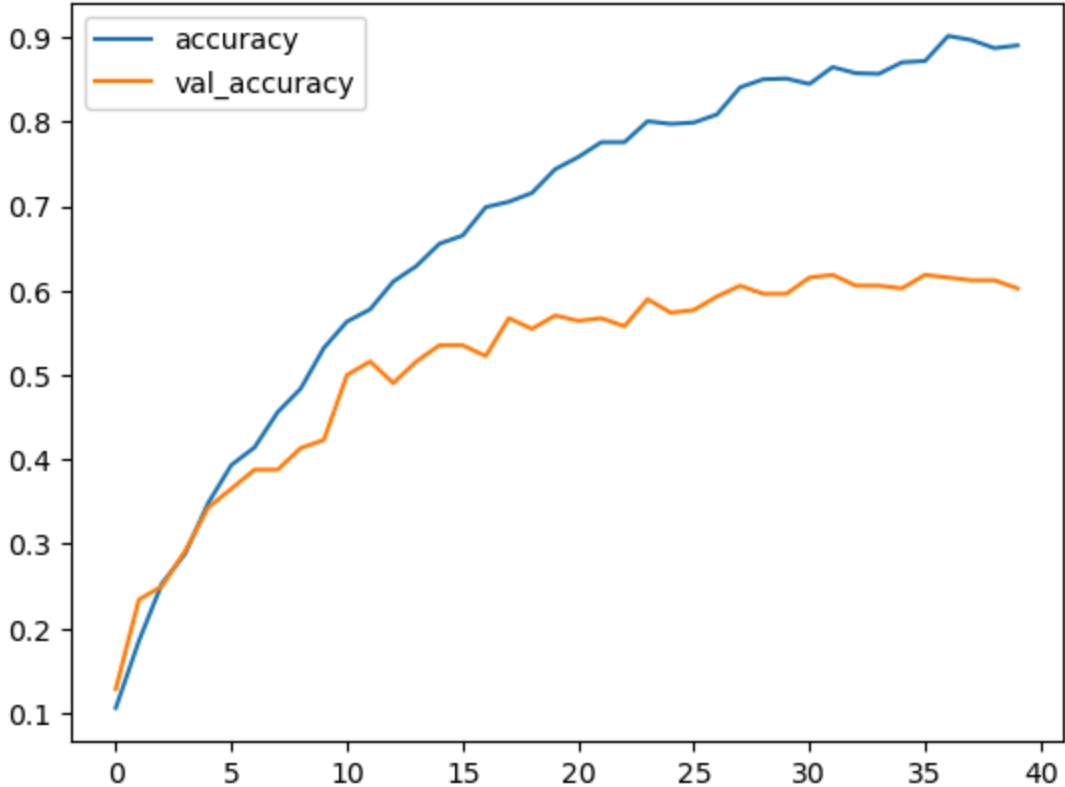
39/39  0s 11ms/step - accuracy: 0.7821 - loss: 0.6458 - val_accuracy: 0.5673 - val_loss: 1.4864
Epoch 23/40

39/39  0s 11ms/step - accuracy: 0.7876 - loss: 0.6228 - val_accuracy: 0.5577 - val_loss: 1.5122
Epoch 24/40

39/39  0s 11ms/step - accuracy: 0.8071 - loss: 0.5947 - val_accuracy: 0.5897 - val_loss: 1.4916
Epoch 25/40

39/39  0s 11ms/step - accuracy: 0.8121 - loss: 0.5334 - val_accuracy: 0.5897 - val_loss: 1.4916

racy: 0.5737 - val_loss: 1.4567
Epoch 26/40
39/39 ————— 0s 11ms/step - accuracy: 0.8067 - loss: 0.5962 - val_accu
racy: 0.5769 - val_loss: 1.5144
Epoch 27/40
39/39 ————— 0s 11ms/step - accuracy: 0.8231 - loss: 0.5233 - val_accu
racy: 0.5929 - val_loss: 1.4102
Epoch 28/40
39/39 ————— 0s 11ms/step - accuracy: 0.8626 - loss: 0.4512 - val_accu
racy: 0.6058 - val_loss: 1.4368
Epoch 29/40
39/39 ————— 0s 11ms/step - accuracy: 0.8385 - loss: 0.4827 - val_accu
racy: 0.5962 - val_loss: 1.4916
Epoch 30/40
39/39 ————— 0s 11ms/step - accuracy: 0.8575 - loss: 0.4516 - val_accu
racy: 0.5962 - val_loss: 1.4568
Epoch 31/40
39/39 ————— 0s 11ms/step - accuracy: 0.8462 - loss: 0.4438 - val_accu
racy: 0.6154 - val_loss: 1.5072
Epoch 32/40
39/39 ————— 0s 11ms/step - accuracy: 0.8613 - loss: 0.4013 - val_accu
racy: 0.6186 - val_loss: 1.4857
Epoch 33/40
39/39 ————— 0s 11ms/step - accuracy: 0.8660 - loss: 0.4079 - val_accu
racy: 0.6058 - val_loss: 1.4646
Epoch 34/40
39/39 ————— 0s 11ms/step - accuracy: 0.8597 - loss: 0.4044 - val_accu
racy: 0.6058 - val_loss: 1.4050
Epoch 35/40
39/39 ————— 0s 11ms/step - accuracy: 0.8760 - loss: 0.3750 - val_accu
racy: 0.6026 - val_loss: 1.4292
Epoch 36/40
39/39 ————— 0s 11ms/step - accuracy: 0.8740 - loss: 0.3570 - val_accu
racy: 0.6186 - val_loss: 1.5568
Epoch 37/40
39/39 ————— 0s 11ms/step - accuracy: 0.9164 - loss: 0.2996 - val_accu
racy: 0.6154 - val_loss: 1.4178
Epoch 38/40
39/39 ————— 0s 11ms/step - accuracy: 0.8913 - loss: 0.3462 - val_accu
racy: 0.6122 - val_loss: 1.5912
Epoch 39/40
39/39 ————— 0s 11ms/step - accuracy: 0.8914 - loss: 0.3233 - val_accu
racy: 0.6122 - val_loss: 1.5286
Epoch 40/40
39/39 ————— 0s 11ms/step - accuracy: 0.8896 - loss: 0.3145 - val_accu
racy: 0.6026 - val_loss: 1.5663



10/10 ————— 0s 4ms/step - accuracy: 0.6046 - loss: 1.5308
Validation accuracy: 0.6026 - Validation loss: 1.5663

```
In [6]: print(f"Validation accuracy: {acc:.4f} - Validation loss: {loss:.4f}")
```

Validation accuracy: 0.6026 - Validation loss: 1.5663

First step complete

mit val_accuracy von über 0.6 und val_loss etwa 1.4

The initial model ($3 \times [\text{Conv+Pool}] + \text{Flatten} + \text{Dense}$) achieved a validation accuracy (**val_acc**) of approximately **60.3%** with a loss (**val_loss**) of **1.57** after 40 epochs. This served as the baseline for subsequent optimizations.

Second step

In this step some of the parameters and hyperparameters are tried to find the best model architecture and training process

- For Modell-Architektur
 - filter number
 - kernel size
 - dense units
 - dropout percentage
- For Training-Prozess
 - weight decay
 - different optimizer
 - lr
 - batch size

```
In [5]: # imports
import keras_tuner as kt
from keras.callbacks import EarlyStopping, TensorBoard
from tensorboard.plugins.hparams import api as hp # TensorBoard HParams
```

```
In [5]: data_augmentation = keras.Sequential([
    layers.RandomFlip("horizontal"),
    layers.RandomRotation(0.15),
    layers.RandomZoom(0.1),
])
```

```
In [ ]: import keras_tuner as kt
from tensorflow import keras
from keras import layers, models, regularizers
import pandas as pd
import numpy as np

# ----- build function -----
def build_full(hp):
    """fix 3xConv+Pool architecture do hp search for filter number / kernel_size /
    Dense units / Dropout / weight-decay / optimizer etc.
    """
    img_height, img_width = 224, 224

    # --- layer hyperparameter -----
    f0 = hp.Choice("f0", [32, 64, 128])
    f1 = hp.Choice("f1", [64, 128, 256])
    f2 = hp.Choice("f2", [128, 256])

    k = hp.Choice("kernel_size", [3, 5, 7])
    dense_units = hp.Choice("dense_units", [64, 128, 256])
    dense_dp = hp.Float("dense_dp", 0.2, 0.5, step=0.1)

    # weight-decay regularization
    wd = hp.Float("weight_decay",
```

```

        min_value=1e-6, max_value=1e-4, sampling="log")
l2_reg = regularizers.l2(wd)

# — the model —————
model = models.Sequential([
    layers.Rescaling(1./255, input_shape=(img_height, img_width, 3)),

    layers.Conv2D(f0, k, padding="same", activation="relu",
                  kernel_regularizer=l2_reg),
    layers.MaxPooling2D(),

    layers.Conv2D(f1, k, padding="same", activation="relu",
                  kernel_regularizer=l2_reg),
    layers.MaxPooling2D(),

    layers.Conv2D(f2, k, padding="same", activation="relu",
                  kernel_regularizer=l2_reg),
    layers.MaxPooling2D(),

    layers.Flatten(),
    layers.Dense(dense_units, activation="relu", kernel_regularizer=l2_reg),
    layers.Dropout(dense_dp),
    layers.Dense(num_classes, activation="softmax")
])

# — train the optimizer —————
opt_name = hp.Choice("optimizer", ["adam", "rmsprop", "sgd"])
lr        = hp.Choice("lr", [1e-3, 5e-4, 1e-4])
if opt_name == "adam":
    optimizer = keras.optimizers.Adam(lr)
elif opt_name == "rmsprop":
    optimizer = keras.optimizers.RMSprop(lr)
else:
    optimizer = keras.optimizers.SGD(lr, momentum=0.9, nesterov=True)

model.compile(optimizer=optimizer,
              loss="sparse_categorical_crossentropy",
              metrics=["accuracy"])

return model

# ----- use batch_size / plateau as hp -----
class FullTuner(kt.tuners.BayesianOptimization):
    def run_trial(self, trial, *fit_args, **fit_kwargs):
        hp = trial.hyperparameters
        # batch size
        fit_kwargs["batch_size"] = hp.Choice("batch_size", [16, 32, 64])

        # early stopping & Plateau
        callbacks = [
            keras.callbacks.EarlyStopping(
                monitor="val_accuracy",
                mode="max",
                patience=10,
                min_delta=0.01,
                restore_best_weights=True)
        ]
        if hp.Boolean("use_plateau"):
            callbacks.append(
                keras.callbacks.ReduceLROnPlateau(
                    monitor="val_loss", factor=0.5, patience=3)
            )
        fit_kwargs["callbacks"] = callbacks

        # super() return history.history (dict) , which satisfy Keras-Tuner require
        return super().run_trial(trial, *fit_args, **fit_kwargs)

# ----- start to search -----
tuner = FullTuner(
    build_full,
    objective="val_accuracy",
    max_trials=40,
    max_consecutive_failed_trials=10,

```



```

        directory="tune_all",
        project_name="cnn_full",
        overwrite=True,
    )

    tuner.search(
        train_ds,
        validation_data=val_ds,
        epochs=40,
    )

    # ----- get the best model -----
    best_hp      = tuner.get_best_hyperparameters(1)[0]
    best_model   = tuner.get_best_models(1)[0]      # get the best model with trained we
    print("★ Best HPs:", best_hp.values)

    val_loss, val_acc = best_model.evaluate(val_ds, verbose=0)
    print(f"★ Best model  val_acc={val_acc:.3f}  val_loss={val_loss:.3f}")

    # ----- trial → CSV -----
    def safe_last(trial, name):
        return trial.metrics.get_last_value(name) if trial.metrics.exists(name) else np

    records = [{
        "trial_id": t.trial_id, **t.hyperparameters.values,
        "status":    t.status,
        "best_epoch": t.best_step,
        "train_acc": safe_last(t, "accuracy"),
        "val_acc":   safe_last(t, "val_accuracy"),
    } for t in tuner.oracle.trials.values()]

    pd.DataFrame(records).to_csv("tuner_results.csv", index=False)
    print("CSV saved → tuner_results.csv")

```

Trial 40 Complete [00h 00m 21s]
 val_accuracy: 0.6121794581413269

Best val_accuracy So Far: 0.6570512652397156
 Total elapsed time: 00h 28m 32s

/home/cxy_otto/miniconda/envs/tf2-stable/lib/python3.11/site-packages/keras/src/saving/saving_lib.py:802: UserWarning: Skipping variable loading for optimizer 'adam', because it has 2 variables whereas the saved optimizer has 22 variables.

```

    saveable.load_own_variables(weights_store.get(inner_path))
★ Best HPs: {'f0': 128, 'f1': 256, 'f2': 256, 'kernel_size': 7, 'dense_units': 128,
'dense_dp': 0.30000000000000004, 'weight_decay': 4.726311111933952e-05, 'optimizer':
'adam', 'lr': 0.0001, 'batch_size': 16, 'use_plateau': True}
★ Best model  val_acc=0.657  val_loss=1.902
CSV saved → tuner_results.csv

```

Architecture and Training Optimization:

Durch Optimierung der Architektur- und Trainingsparameter (Anzahl der Filter, Kernelgröße, Dropout-Rate, Gewicht-Regularisierung, Lernrate, Batch-Größe und Lernraten-Scheduler) wurde ein deutlicher Fortschritt erzielt. Die besten gefundenen Hyperparameter (f0=128, f1=256, f2=256, Kernelgröße=7, Dense=128, Dropout=0,3, Gewicht-Regularisierung=4,73e-5, Adam-Optimizer mit LR=1e-4 und Batchgröße=16) führten zu einer signifikanten Steigerung der Validierungsgenauigkeit auf etwa **65,7 %** bei einem Verlust von **1,90**. Die größere Kernelgröße (7) und eine stärkere Modellkapazität trugen dabei entscheidend zur Verbesserung bei.

Optimization of architectural and training parameters — including filter sizes, kernel size, dropout rates, weight decay, learning rate, batch size, and the learning rate scheduler—resulted in significant improvements. The optimal hyperparameters (f0=128, f1=256, f2=256, kernel size=7, dense units=128, dropout=0.3, weight decay=4.73e-5, Adam optimizer with LR=1e-4, and batch size=16) notably improved validation accuracy to about **65.7%**, although with a slightly higher

loss (1.90). Larger kernel sizes (7) and increased model capacity played a critical role in this improvement.

Step 2.2 Data Augmentation

Because of the risk of exploding video memory, I separated this out and trained it on the optimal model some methods are used here

- flip(horizontal flip)
- rot RandomRotation
- zoom RandomZoom

```
In [ ]: # ----- fix the best parameters from last search -----
BEST = {
    "f0": 128, "f1": 256, "f2": 256,      # three Conv blocks
    "k" : 7,                             # kernel_size
    "dense": 128, "dense_dp": 0.3,
    "wd": 4.726311111933952e-05,         # weight-decay
    "lr": 1e-4, "batch_size": 16,
    "use_plateau": True                 # use ReduceLROnPlateau callback
}
l2_reg = regularizers.l2(BEST["wd"])

# ----- search for data augmentation HyperModel -----
def build_aug_model(hp):
    # the data augmentation search space
    rot = hp.Float("rot", 0.0, 0.25, step=0.05)  # 0 → 0.25 (≈ ±45°)
    zoom = hp.Float("zoom", 0.0, 0.25, step=0.05)
    flip = hp.Boolean("flip")

    aug = keras.Sequential([
        layers.Rescaling(1./255, input_shape=(224, 224, 3)),
        layers.RandomFlip("horizontal") if flip else layers.Lambda(lambda x:x),
        layers.RandomRotation(rot) if rot>0 else layers.Lambda(lambda x:x),
        layers.RandomZoom(zoom) if zoom>0 else layers.Lambda(lambda x:x),
    ], name="augmentation")

    # fix the BEST parameters of model architecture and training
    m = models.Sequential([
        aug,
        layers.Conv2D(BEST["f0"], BEST["k"], padding="same", activation="relu", kernel_regularizer=l2_reg),
        layers.MaxPooling2D(),
        layers.Conv2D(BEST["f1"], BEST["k"], padding="same", activation="relu", kernel_regularizer=l2_reg),
        layers.MaxPooling2D(),
        layers.Conv2D(BEST["f2"], BEST["k"], padding="same", activation="relu", kernel_regularizer=l2_reg),
        layers.MaxPooling2D(),
        layers.Flatten(),
        layers.Dense(BEST["dense"], activation="relu", kernel_regularizer=l2_reg),
        layers.Dropout(BEST["dense_dp"]),
        layers.Dense(num_classes, activation="softmax")
    ])

    opt = keras.optimizers.Adam(learning_rate=BEST["lr"])
    m.compile(opt, loss="sparse_categorical_crossentropy", metrics=["accuracy"])
    return m

# ----- Tuner search rot/zoom/flip -----
tuner = kt.BayesianOptimization(
    build_aug_model,
    objective="val_accuracy",
    max_trials=15,
    directory="aug_search",
    project_name="cnn_best_aug",
    overwrite=True
)

# callback (same as BEST model)
```

```

callbacks = [
    keras.callbacks.EarlyStopping(monitor="val_accuracy", mode="max",
                                  patience=8, min_delta=0.01,
                                  restore_best_weights=True)
]
if BEST["use_plateau"]:
    callbacks.append(
        keras.callbacks.ReduceLROnPlateau(monitor="val_loss",
                                            factor=0.5, patience=3)
    )

tuner.search(
    train_ds,
    validation_data=val_ds,
    epochs=40,
    batch_size=BEST["batch_size"],  # ← fix batch size
    callbacks=callbacks
)

# ----- read & evaluation for the final model -----
best_aug_hp = tuner.get_best_hyperparameters(1)[0]
print("★ Best augmentation HPs:", best_aug_hp.values)

final_model = tuner.get_best_models(1)[0]
val_loss, val_acc = final_model.evaluate(val_ds, verbose=0)
print(f"★ Final model  val_acc={val_acc:.3f} | val_loss={val_loss:.3f}")

```

Trial 15 Complete [00h 02m 04s]
val_accuracy: 0.7564102411270142

Best val_accuracy So Far: 0.7916666865348816

Total elapsed time: 00h 34m 19s

★ Best augmentation HPs: {'rot': 0.1, 'zoom': 0.0, 'flip': False}

/home/cxy_otto/miniconda/envs/tf2-stable/lib/python3.11/site-packages/keras/src/saving/saving_lib.py:802: UserWarning: Skipping variable loading for optimizer 'adam', because it has 2 variables whereas the saved optimizer has 22 variables.
 saveable.load_own_variables(weights_store.get(inner_path))

★ Final model val_acc=0.792 | val_loss=1.047

Optimization via Data Augmentation:

Additionally, a systematic Bayesian optimization of data augmentation parameters (rotation, zoom, horizontal flip) was conducted. The optimal settings (rotation: 10%, zoom: 0%, flip: False) further significantly increased validation accuracy to around **79.2%** while simultaneously reducing loss to **1.05**. Moderate rotation proved particularly beneficial in enhancing the robustness of the model.

Step 3. Visualization of the training prozess of the best model

in this step, all the things are showed in tensorboard.

In [3]: `import tensorflow as tf, numpy as np, matplotlib.pyplot as plt, io, pathlib, datetime, os`
`from sklearn.metrics import confusion_matrix, classification_report`
`from tensorflow.keras import layers, models, regularizers`

In [14]: `import tensorflow as tf, numpy as np, matplotlib.pyplot as plt`
`from tensorflow.keras import layers, models, regularizers`
`from sklearn.metrics import confusion_matrix, classification_report`
`import pathlib, datetime, io, os`

```

# ----- (2) fix the best HP -----
BEST = {
    "f0": 128, "f1": 256, "f2": 256, "k": 7,
    "dense": 128, "dense_dp": 0.3,
    "wd": 4.726311111933952e-05,
    "lr": 1e-4, "batch": 16, "use_plateau": True,
    "rot": 0.10, "zoom": 0.0, "flip": False,

```

```

}

# ----- (3) final -----
def build_final():
    aug = tf.keras.Sequential([layers.Rescaling(1/255.)], name="aug")
    if BEST["flip"]: aug.add(layers.RandomFlip("horizontal"))
    if BEST["rot"] : aug.add(layers.RandomRotation(BEST["rot"]))
    if BEST["zoom"]: aug.add(layers.RandomZoom(BEST["zoom"]))

    l2 = regularizers.l2(BEST["wd"])
    model = models.Sequential([
        aug,
        layers.Conv2D(BEST["f0"], BEST["k"], activation="relu",
                      padding="same", kernel_regularizer=l2),
        layers.MaxPooling2D(),
        layers.Conv2D(BEST["f1"], BEST["k"], activation="relu",
                      padding="same", kernel_regularizer=l2),
        layers.MaxPooling2D(),
        layers.Conv2D(BEST["f2"], BEST["k"], activation="relu",
                      padding="same", kernel_regularizer=l2),
        layers.MaxPooling2D(),
        layers.Flatten(),
        layers.Dense(BEST["dense"], activation="relu", kernel_regularizer=l2),
        layers.Dropout(BEST["dense_dp"]),
        layers.Dense(num_classes, activation="softmax")
    ])
    model.compile(
        optimizer=tf.keras.optimizers.Adam(BEST["lr"]),
        loss="sparse_categorical_crossentropy",
        metrics=["accuracy"])
    return model

model = build_final()

# ----- (4) dir preparation -----
stamp = datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
tb_dir = pathlib.Path("tb_final", stamp)
art_dir = pathlib.Path("eval_artifacts"); art_dir.mkdir(parents=True, exist_ok=True)

# ----- (5) training + TensorBoard -----
cbs = [
    tf.keras.callbacks.TensorBoard(log_dir=tb_dir, update_freq="epoch"),
    tf.keras.callbacks.EarlyStopping(
        monitor="val_accuracy", patience=10, min_delta=0.01, restore_best_weights=True)
]
if BEST["use_plateau"]:
    cbs.append(
        tf.keras.callbacks.ReduceLROnPlateau(
            monitor="val_loss", factor=0.5, patience=3)
    )

history = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=50,
    batch_size=BEST["batch"],
    callbacks=cbs
)

# ----- (6) confusion matrix -----
y_true = np.concatenate([y for _, y in val_ds], axis=0)
y_pred = model.predict(val_ds, verbose=0).argmax(axis=1)

cm = confusion_matrix(y_true, y_pred)
fig, ax = plt.subplots(figsize=(6,6))
im = ax.imshow(cm, cmap="Blues")
ax.set_xticks(range(num_classes), class_names, rotation=45, ha="right")
ax.set_yticks(range(num_classes), class_names)
for (i,j), v in np.ndenumerate(cm):
    ax.text(j, i, int(v), ha='center', va='center',
            color="white" if v > cm.max()*0.6 else "black", fontsize=7)
plt.colorbar(im, fraction=0.046)
ax.set_xlabel("Predicted"); ax.set_ylabel("True"); plt.tight_layout()

```

```

cm_png = art_dir / "confusion_matrix.png"; fig.savefig(cm_png, dpi=300); plt.close()

# TensorBoard -> Images
with tf.summary.create_file_writer(str(tb_dir)).as_default():
    png_bytes = tf.io.read_file(str(cm_png))
    img_tensor = tf.io.decode_png(png_bytes, channels=4)[tf.newaxis]
    tf.summary.image("confusion_matrix", img_tensor, step=0)

report = classification_report(y_true, y_pred, target_names=class_names, digits=3)
(art_dir / "classification_report.txt").write_text(report)
print(report)

# ----- (7) training prozess PNG -----
plt.figure(figsize=(6,4))
plt.plot(history.history["accuracy"], label="train_acc")
plt.plot(history.history["val_accuracy"], label="val_acc")
plt.plot(history.history["loss"], label="train_loss", ls="--")
plt.plot(history.history["val_loss"], label="val_loss", ls="--")
plt.xlabel("Epoch"); plt.ylabel("Metric"); plt.legend(); plt.tight_layout()
curve_png = art_dir / "training_curves.png"; plt.savefig(curve_png, dpi=300); plt.c

print(f"\ndone!\nTensorBoard logs : {tb_dir}\n"
      f"Confusion matrix : {cm_png}\n"
      f"Classification : {art_dir/'classification_report.txt'}\n"
      f"Curves : {curve_png}\n")
print("\nrn: tensorboard serve --logdir tb_final --port 6006")

# ----- (8) In-Notebook Training Curve Visualization -----
import matplotlib.pyplot as plt

# Plot accuracy curves
plt.figure(figsize=(8, 4))
plt.plot(history.history["accuracy"], label="Train Accuracy")
plt.plot(history.history["val_accuracy"], label="Val Accuracy")
plt.title("Training vs. Validation Accuracy")
plt.xlabel("Epoch")
plt.ylabel("Accuracy")
plt.legend()
plt.grid(alpha=0.3)
plt.tight_layout()
plt.show()

# Plot Loss curves
plt.figure(figsize=(8, 4))
plt.plot(history.history["loss"], label="Train Loss", linestyle="--")
plt.plot(history.history["val_loss"], label="Val Loss", linestyle="--")
plt.title("Training vs. Validation Loss")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.legend()
plt.grid(alpha=0.3)
plt.tight_layout()
plt.show()

```

Epoch 1/50

39/39 ————— 5s 102ms/step - accuracy: 0.0880 - loss: 2.3712 - val_accuracy: 0.0833 - val_loss: 2.3288 - learning_rate: 1.0000e-04
Epoch 2/50

39/39 ————— 4s 97ms/step - accuracy: 0.1258 - loss: 2.2979 - val_accuracy: 0.2468 - val_loss: 2.1412 - learning_rate: 1.0000e-04
Epoch 3/50

39/39 ————— 5s 97ms/step - accuracy: 0.2405 - loss: 2.1019 - val_accuracy: 0.2276 - val_loss: 2.1312 - learning_rate: 1.0000e-04
Epoch 4/50

39/39 ————— 4s 98ms/step - accuracy: 0.2809 - loss: 1.9760 - val_accuracy: 0.2885 - val_loss: 1.9675 - learning_rate: 1.0000e-04
Epoch 5/50

39/39 ————— 4s 97ms/step - accuracy: 0.3324 - loss: 1.8671 - val_accuracy: 0.3526 - val_loss: 1.7969 - learning_rate: 1.0000e-04
Epoch 6/50

39/39 ————— 4s 98ms/step - accuracy: 0.3487 - loss: 1.7191 - val_accuracy: 0.4455 - val_loss: 1.6624 - learning_rate: 1.0000e-04
Epoch 7/50

39/39 ————— 4s 98ms/step - accuracy: 0.4444 - loss: 1.5573 - val_accuracy: 0.4968 - val_loss: 1.4886 - learning_rate: 1.0000e-04
Epoch 8/50

39/39 ————— 4s 99ms/step - accuracy: 0.5075 - loss: 1.4020 - val_accuracy: 0.5481 - val_loss: 1.4585 - learning_rate: 1.0000e-04
Epoch 9/50

39/39 ————— 4s 99ms/step - accuracy: 0.5437 - loss: 1.3018 - val_accuracy: 0.6154 - val_loss: 1.3170 - learning_rate: 1.0000e-04
Epoch 10/50

39/39 ————— 4s 98ms/step - accuracy: 0.5849 - loss: 1.1981 - val_accuracy: 0.5032 - val_loss: 1.8169 - learning_rate: 1.0000e-04
Epoch 11/50

39/39 ————— 4s 98ms/step - accuracy: 0.5990 - loss: 1.1676 - val_accuracy: 0.5641 - val_loss: 1.5308 - learning_rate: 1.0000e-04
Epoch 12/50

39/39 ————— 4s 99ms/step - accuracy: 0.6121 - loss: 1.0927 - val_accuracy: 0.6410 - val_loss: 1.1143 - learning_rate: 1.0000e-04
Epoch 13/50

39/39 ————— 4s 98ms/step - accuracy: 0.6746 - loss: 0.9464 - val_accuracy: 0.6538 - val_loss: 1.1685 - learning_rate: 1.0000e-04
Epoch 14/50

39/39 ————— 4s 98ms/step - accuracy: 0.7060 - loss: 0.8684 - val_accuracy: 0.6635 - val_loss: 1.1586 - learning_rate: 1.0000e-04
Epoch 15/50

39/39 ————— 4s 99ms/step - accuracy: 0.7299 - loss: 0.8428 - val_accuracy: 0.6795 - val_loss: 1.1220 - learning_rate: 1.0000e-04
Epoch 16/50

39/39 ————— 4s 98ms/step - accuracy: 0.7483 - loss: 0.7749 - val_accuracy: 0.6667 - val_loss: 1.0542 - learning_rate: 5.0000e-05
Epoch 17/50

39/39 ————— 4s 99ms/step - accuracy: 0.7820 - loss: 0.7196 - val_accuracy: 0.6923 - val_loss: 1.1648 - learning_rate: 5.0000e-05
Epoch 18/50

39/39 ————— 4s 98ms/step - accuracy: 0.7912 - loss: 0.6646 - val_accuracy: 0.7019 - val_loss: 1.0151 - learning_rate: 5.0000e-05
Epoch 19/50

39/39 ————— 4s 98ms/step - accuracy: 0.7749 - loss: 0.6898 - val_accuracy: 0.6955 - val_loss: 1.2185 - learning_rate: 5.0000e-05
Epoch 20/50

39/39 ————— 4s 98ms/step - accuracy: 0.7776 - loss: 0.6392 - val_accuracy: 0.7019 - val_loss: 1.3423 - learning_rate: 5.0000e-05
Epoch 21/50

39/39 ————— 4s 99ms/step - accuracy: 0.8103 - loss: 0.6131 - val_accuracy: 0.7212 - val_loss: 1.0653 - learning_rate: 5.0000e-05
Epoch 22/50

39/39 ————— 4s 98ms/step - accuracy: 0.8153 - loss: 0.5805 - val_accuracy: 0.7308 - val_loss: 1.0485 - learning_rate: 2.5000e-05
Epoch 23/50

39/39 ————— 4s 98ms/step - accuracy: 0.8424 - loss: 0.5234 - val_accuracy: 0.7404 - val_loss: 0.9707 - learning_rate: 2.5000e-05
Epoch 24/50

39/39 ————— 4s 98ms/step - accuracy: 0.8402 - loss: 0.5180 - val_accuracy: 0.7276 - val_loss: 1.0582 - learning_rate: 2.5000e-05
Epoch 25/50

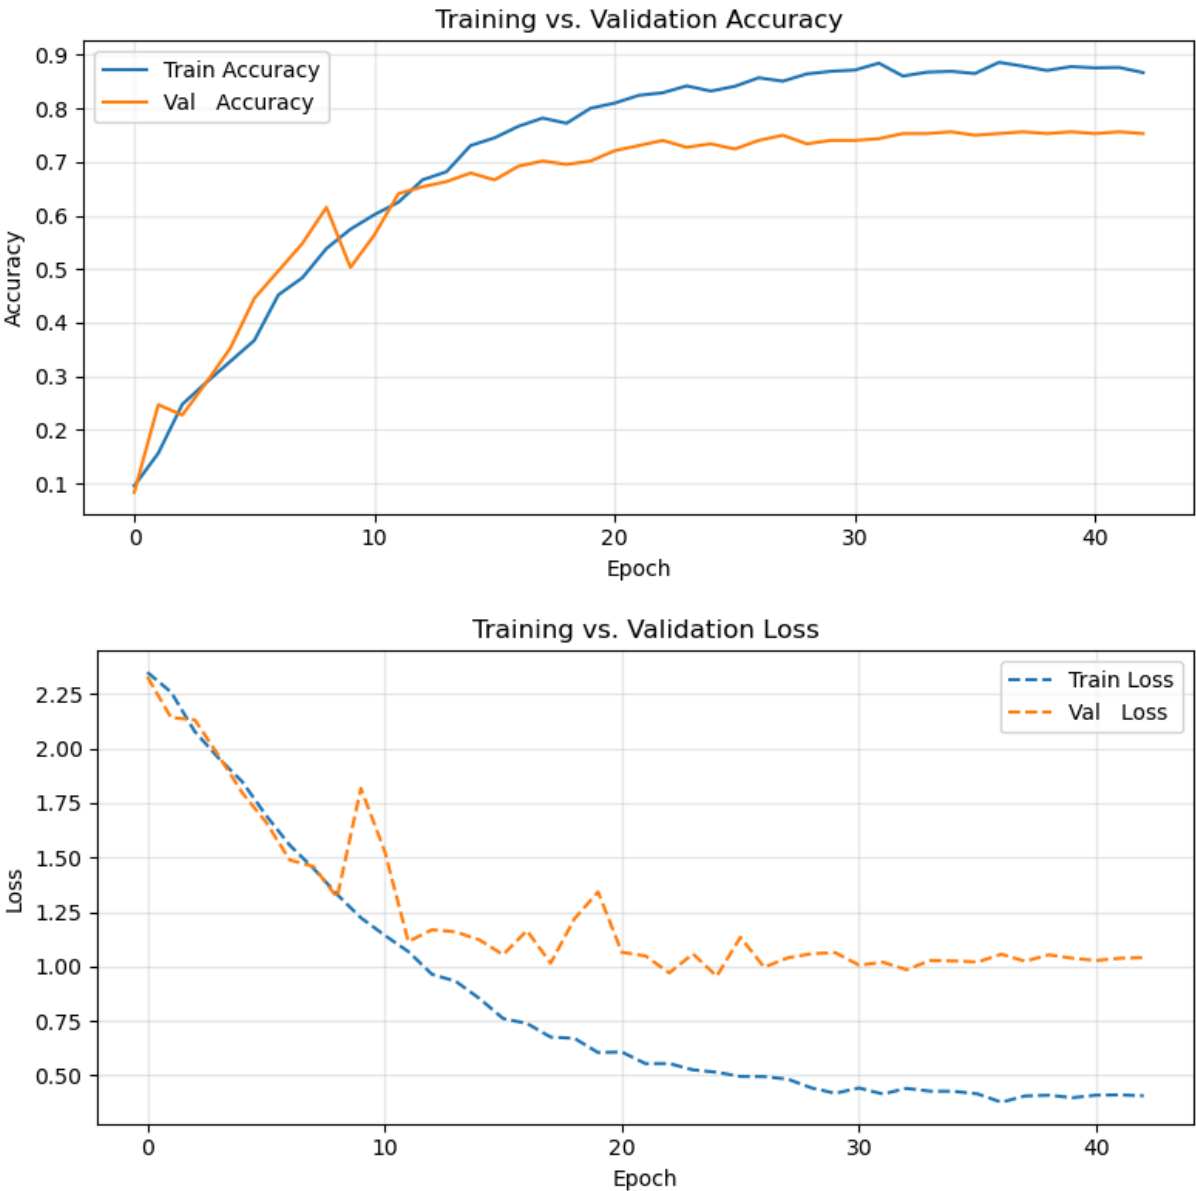
39/39 ————— 4s 98ms/step - accuracy: 0.8461 - loss: 0.4883 - val_accuracy:

racy: 0.7340 - val_loss: 0.9554 - learning_rate: 2.5000e-05
Epoch 26/50
39/39 ————— 4s 98ms/step - accuracy: 0.8406 - loss: 0.4978 - val_accu
racy: 0.7244 - val_loss: 1.1341 - learning_rate: 2.5000e-05
Epoch 27/50
39/39 ————— 4s 99ms/step - accuracy: 0.8597 - loss: 0.4811 - val_accu
racy: 0.7404 - val_loss: 0.9968 - learning_rate: 2.5000e-05
Epoch 28/50
39/39 ————— 4s 98ms/step - accuracy: 0.8717 - loss: 0.4343 - val_accu
racy: 0.7500 - val_loss: 1.0400 - learning_rate: 2.5000e-05
Epoch 29/50
39/39 ————— 4s 98ms/step - accuracy: 0.8756 - loss: 0.4252 - val_accu
racy: 0.7340 - val_loss: 1.0586 - learning_rate: 1.2500e-05
Epoch 30/50
39/39 ————— 4s 98ms/step - accuracy: 0.8700 - loss: 0.4064 - val_accu
racy: 0.7404 - val_loss: 1.0636 - learning_rate: 1.2500e-05
Epoch 31/50
39/39 ————— 4s 98ms/step - accuracy: 0.8794 - loss: 0.4197 - val_accu
racy: 0.7404 - val_loss: 1.0062 - learning_rate: 1.2500e-05
Epoch 32/50
39/39 ————— 4s 97ms/step - accuracy: 0.8778 - loss: 0.4197 - val_accu
racy: 0.7436 - val_loss: 1.0199 - learning_rate: 6.2500e-06
Epoch 33/50
39/39 ————— 4s 98ms/step - accuracy: 0.8750 - loss: 0.4105 - val_accu
racy: 0.7532 - val_loss: 0.9855 - learning_rate: 6.2500e-06
Epoch 34/50
39/39 ————— 4s 97ms/step - accuracy: 0.8651 - loss: 0.4485 - val_accu
racy: 0.7532 - val_loss: 1.0277 - learning_rate: 6.2500e-06
Epoch 35/50
39/39 ————— 4s 97ms/step - accuracy: 0.8709 - loss: 0.4488 - val_accu
racy: 0.7564 - val_loss: 1.0256 - learning_rate: 3.1250e-06
Epoch 36/50
39/39 ————— 4s 97ms/step - accuracy: 0.8767 - loss: 0.3833 - val_accu
racy: 0.7500 - val_loss: 1.0207 - learning_rate: 3.1250e-06
Epoch 37/50
39/39 ————— 4s 97ms/step - accuracy: 0.8987 - loss: 0.3563 - val_accu
racy: 0.7532 - val_loss: 1.0565 - learning_rate: 3.1250e-06
Epoch 38/50
39/39 ————— 4s 97ms/step - accuracy: 0.8700 - loss: 0.4160 - val_accu
racy: 0.7564 - val_loss: 1.0242 - learning_rate: 1.5625e-06
Epoch 39/50
39/39 ————— 4s 97ms/step - accuracy: 0.8791 - loss: 0.4067 - val_accu
racy: 0.7532 - val_loss: 1.0536 - learning_rate: 1.5625e-06
Epoch 40/50
39/39 ————— 4s 98ms/step - accuracy: 0.8770 - loss: 0.3943 - val_accu
racy: 0.7564 - val_loss: 1.0377 - learning_rate: 1.5625e-06
Epoch 41/50
39/39 ————— 4s 98ms/step - accuracy: 0.8757 - loss: 0.4209 - val_accu
racy: 0.7532 - val_loss: 1.0276 - learning_rate: 7.8125e-07
Epoch 42/50
39/39 ————— 4s 98ms/step - accuracy: 0.8770 - loss: 0.3919 - val_accu
racy: 0.7564 - val_loss: 1.0380 - learning_rate: 7.8125e-07
Epoch 43/50
39/39 ————— 4s 98ms/step - accuracy: 0.8629 - loss: 0.3889 - val_accu
racy: 0.7532 - val_loss: 1.0406 - learning_rate: 7.8125e-07
2025-07-14 19:20:37.795831: W tensorflow/core/framework/local_rendezvous.cc:404] Loc
al rendezvous is aborting with status: OUT_OF_RANGE: End of sequence

	precision	recall	f1-score	support
Abstandshalter	0.963	0.963	0.963	27
Auslassventil	0.923	0.800	0.857	30
Blechlineal	0.727	0.593	0.653	27
Filterkartusche	0.964	1.000	0.982	27
Gewindestange	0.800	0.800	0.800	35
Hohlschraube	0.667	0.643	0.655	28
Hutmutter	0.759	0.667	0.710	33
Hydraulikstutzen	0.667	0.690	0.678	29
Nutenstein	0.588	0.541	0.563	37
Schraubenfeder	0.630	0.872	0.731	39
accuracy			0.753	312
macro avg	0.769	0.757	0.759	312
weighted avg	0.759	0.753	0.752	312

done!
TensorBoard logs : tb_final/20250714-191750
Confusion matrix : eval_artifacts/confusion_matrix.png
Classification : eval_artifacts/classification_report.txt
Curves : eval_artifacts/training_curves.png

run: tensorboard serve --logdir tb_final --port 6006



This cell is for visualization of everything with Tensorboard.

```
In [17]: ! tensorboard serve --logdir tb_final --port 6006
```



```
/home/cxy_otto/miniconda/envs/tf2-stable/lib/python3.11/site-packages/tensorboard/default.py:30: UserWarning: pkg_resources is deprecated as an API. See https://setuptools.pypa.io/en/latest/pkg_resources.html. The pkg_resources package is slated for removal as early as 2025-11-30. Refrain from using this package or pin to Setuptools<81.
```

```
import pkg_resources
2025-07-14 19:32:51.972276: I tensorflow/core/util/port.cc:113] oneDNN custom operations are on. You may see slightly different numerical results due to floating-point round-off errors from different computation orders. To turn them off, set the environment variable `TF_ENABLE_ONEDNN_OPTS=0`.
2025-07-14 19:32:51.998896: I tensorflow/core/platform/cpu_feature_guard.cc:210] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.
To enable the following instructions: SSE4.1 SSE4.2 AVX AVX2 AVX_VNNI FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.
2025-07-14 19:32:53.465873: I external/local_xla/xla/stream_executor/cuda/cuda_executor.cc:984] could not open file to read NUMA node: /sys/bus/pci/devices/0000:01:00.0/numa_node
Your kernel may have been built without NUMA support.
2025-07-14 19:32:53.529404: I external/local_xla/xla/stream_executor/cuda/cuda_executor.cc:984] could not open file to read NUMA node: /sys/bus/pci/devices/0000:01:00.0/numa_node
Your kernel may have been built without NUMA support.
2025-07-14 19:32:53.529467: I external/local_xla/xla/stream_executor/cuda/cuda_executor.cc:984] could not open file to read NUMA node: /sys/bus/pci/devices/0000:01:00.0/numa_node
Your kernel may have been built without NUMA support.
```

NOTE: Using experimental fast data loading logic. To disable, pass
"--load_fast=false" and report issues on GitHub. More details:
<https://github.com/tensorflow/tensorboard/issues/4784>

Serving TensorBoard on localhost; to expose to the network, use a proxy or pass --bind_all
TensorBoard 2.16.2 at http://localhost:6006/ (Press CTRL+C to quit)
^C

evaluation of model showing in the notebook

In [34]: *# 1) Confusion Matrix & Classification Report*

```
# gather ground truth and predictions
y_true = np.concatenate([y for _, y in val_ds], axis=0)
y_pred = model.predict(val_ds, verbose=0).argmax(axis=1)

# compute confusion matrix
cm = confusion_matrix(y_true, y_pred)

# plot heatmap
plt.figure(figsize=(8,8))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues",
            xticklabels=class_names, yticklabels=class_names,
            cbar_kws=dict(label="count"))
plt.xlabel("Predicted label")
plt.ylabel("True label")
plt.title("Confusion Matrix - Final Model")
plt.tight_layout()
plt.show()

# print classification report
print("Classification Report - Final Model\n")
print(classification_report(
    y_true, y_pred, target_names=class_names, digits=3, zero_division=0
))
```

```
2025-07-14 20:07:17.328627: W tensorflow/core/framework/local_rendezvous.cc:404] Local rendezvous is aborting with status: OUT_OF_RANGE: End of sequence
```

TensorBoard

TIME SERIES

SCALARS

Filter runs (regex)

Run

aug/trial_00

aug/trial_01

aug/trial_02

aug/trial_03

aug/trial_04

aug/trial_05

aug/trial_06

aug/trial_07

aug/trial_08

aug/trial_09

aug/trial_10

aug/trial_11

aug/trial_12

aug/trial_13

aug/trial_14

full/trial_00

full/trial_01

full/trial_02

full/trial_03

full/trial_04

full/trial_05

full/trial_06

full/trial_07

Filter tags (regex)

loss

val_accuracy

val_loss

Run

Value

Step

Time

Relative

aug/trial_00 0.7083 33 7/13/25, 11:58 PM 0

aug/trial_01 0.6392 19 7/13/25, 11:58 PM 0

aug/trial_02 0.6859 30 7/13/25, 11:58 PM 0

aug/trial_03 0.7019 19 7/13/25, 11:58 PM 0

aug/trial_04 0.7468 30 7/13/25, 11:58 PM 0

aug/trial_05 0.7756 24 7/13/25, 11:58 PM 0

aug/trial_06 0.6891 38 7/13/25, 11:58 PM 0

aug/trial_07 0.6859 25 7/13/25, 11:58 PM 0

aug/trial_08 0.7756 31 7/13/25, 11:58 PM 0

aug/trial_09 0.7917 39 7/13/25, 11:58 PM 0

aug/trial_10 0.7564 28 7/13/25, 11:58 PM 0

aug/trial_11 0.7692 31 7/13/25, 11:58 PM 0

aug/trial_12 0.7724 29 7/13/25, 11:58 PM 0

aug/trial_13 0.7564 22 7/13/25, 11:58 PM 0

aug/trial_14 0.5897 20 7/13/25, 11:58 PM 0

full/trial_00 0.4295 35 7/13/25, 11:58 PM 0

full/trial_01 0.5897 24 7/13/25, 11:58 PM 0

full/trial_02 0.09295 2 7/13/25, 11:58 PM 0

full/trial_03 0.5962 25 7/13/25, 11:58 PM 0

full/trial_04 0.5962 16 7/13/25, 11:58 PM 0

full/trial_05 0.5962 14 7/13/25, 11:58 PM 0

full/trial_06 0.5962 34 7/13/25, 11:58 PM 0

full/trial_07 0.2115 20 7/13/25, 11:58 PM 0

full/trial_08 0.4295 20 7/13/25, 11:58 PM 0

full/trial_09 0.08654 0 7/13/25, 11:58 PM 0

full/trial_10 0.25 15 7/13/25, 11:58 PM 0

full/trial_11 0.6186 18 7/13/25, 11:58 PM 0

full/trial_12 0.5128 38 7/13/25, 11:58 PM 0

full/trial_13 0.5865 37 7/13/25, 11:58 PM 0

full/trial_14 0.2212 10 7/13/25, 11:58 PM 0

full/trial_15 0.6314 30 7/13/25, 11:58 PM 0

full/trial_16 0.5929 26 7/13/25, 11:58 PM 0

full/trial_17 0.4199 39 7/13/25, 11:58 PM 0

full/trial_18 0.6282 16 7/13/25, 11:58 PM 0

full/trial_19 0.2083 14 7/13/25, 11:58 PM 0

full/trial_20 0.5897 34 7/13/25, 11:58 PM 0

full/trial_21 0.1923 16 7/13/25, 11:58 PM 0

full/trial_22 0.6538 17 7/13/25, 11:58 PM 0

full/trial_23 0.09295 0 7/13/25, 11:58 PM 0

full/trial_24 0.3077 16 7/13/25, 11:58 PM 0

TensorBoard

TIME SERIES

SCALARS

Filter runs (regex)

Run

aug/trial_00

aug/trial_01

aug/trial_02

aug/trial_03

aug/trial_04

aug/trial_05

aug/trial_06

aug/trial_07

aug/trial_08

aug/trial_09

aug/trial_10

aug/trial_11

aug/trial_12

aug/trial_13

aug/trial_14

full/trial_00

full/trial_01

full/trial_02

full/trial_03

full/trial_04

full/trial_05

full/trial_06

full/trial_07

Filter tags (regex)

loss

val_accuracy

val_loss

Run

Value

Step

Time

Relative

aug/trial_00 1.077 33 7/13/25, 11:58 PM 0

aug/trial_01 1.886 19 7/13/25, 11:58 PM 0

aug/trial_02 1.248 30 7/13/25, 11:58 PM 0

aug/trial_03 1.118 19 7/13/25, 11:58 PM 0

aug/trial_04 0.9967 30 7/13/25, 11:58 PM 0

aug/trial_05 1.036 34 7/13/25, 11:58 PM 0

aug/trial_06 1.014 24 7/13/25, 11:58 PM 0

aug/trial_07 1.09 38 7/13/25, 11:58 PM 0

aug/trial_08 2.233 25 7/13/25, 11:58 PM 0

aug/trial_09 1.041 31 7/13/25, 11:58 PM 0

aug/trial_10 1.047 39 7/13/25, 11:58 PM 0

aug/trial_11 1.156 28 7/13/25, 11:58 PM 0

aug/trial_12 1.056 31 7/13/25, 11:58 PM 0

aug/trial_13 0.8925 29 7/13/25, 11:58 PM 0

aug/trial_14 1.085 22 7/13/25, 11:58 PM 0

full/trial_00 2.167 20 7/13/25, 11:58 PM 0

full/trial_01 1.776 35 7/13/25, 11:58 PM 0

full/trial_02 2.446 24 7/13/25, 11:58 PM 0

full/trial_03 2.315 2 7/13/25, 11:58 PM 0

full/trial_04 1.542 25 7/13/25, 11:58 PM 0

full/trial_05 1.902 16 7/13/25, 11:58 PM 0

full/trial_06 3.006 14 7/13/25, 11:58 PM 0

full/trial_07 1.787 34 7/13/25, 11:58 PM 0

full/trial_08 2.282 20 7/13/25, 11:58 PM 0

full/trial_09 6.592 20 7/13/25, 11:58 PM 0

full/trial_10 2.306 0 7/13/25, 11:58 PM 0

full/trial_11 2.141 15 7/13/25, 11:58 PM 0

full/trial_12 1.684 18 7/13/25, 11:58 PM 0

full/trial_13 1.68 38 7/13/25, 11:58 PM 0

full/trial_14 1.41 37 7/13/25, 11:58 PM 0

full/trial_15 2.267 10 7/13/25, 11:58 PM 0

full/trial_16 2.627 30 7/13/25, 11:58 PM 0

full/trial_17 2.363 26 7/13/25, 11:58 PM 0

full/trial_18 1.67 39 7/13/25, 11:58 PM 0

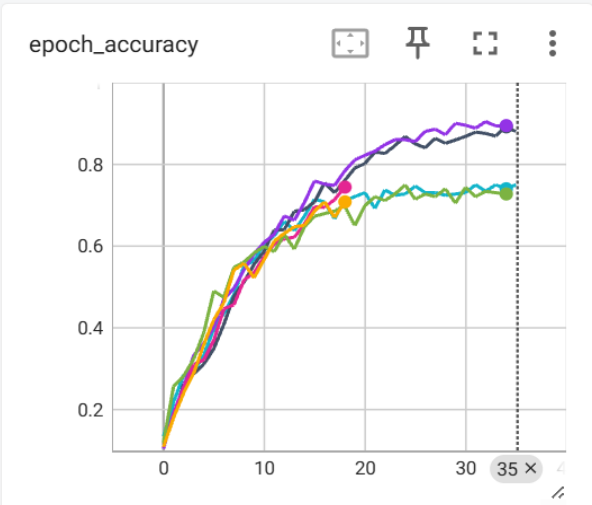
full/trial_19 1.953 16 7/13/25, 11:58 PM 0

full/trial_20 2.268 14 7/13/25, 11:58 PM 0

full/trial_21 1.558 34 7/13/25, 11:58 PM 0

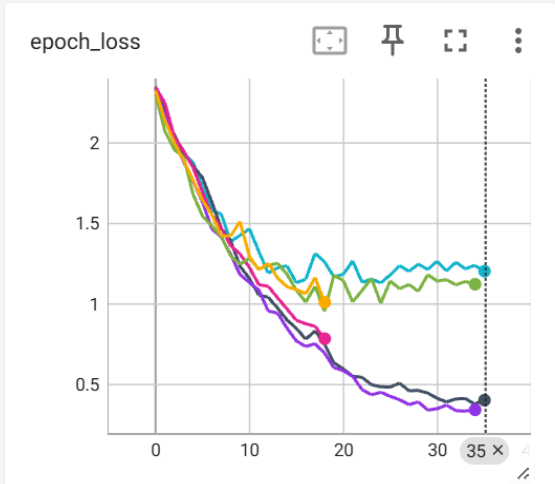
full/trial_22 2.282 16 7/13/25, 11:58 PM 0

epoch_accuracy



Run	Value	Step	Time	Relative
20250714-001210/train	0.8918	34	7/14/25, 12:14 AM	2.075 min
20250714-001210/validation	0.7404	34	7/14/25, 12:14 AM	2.075 min
20250714-002011/train	0.7444	18	7/14/25, 12:21 AM	1.137 min
20250714-002011/validation	0.7083	18	7/14/25, 12:21 AM	1.137 min
20250714-002131/train	0.895	34	7/14/25, 12:23 AM	2.131 min
20250714-002131/validation	0.7276	34	7/14/25, 12:23 AM	2.131 min

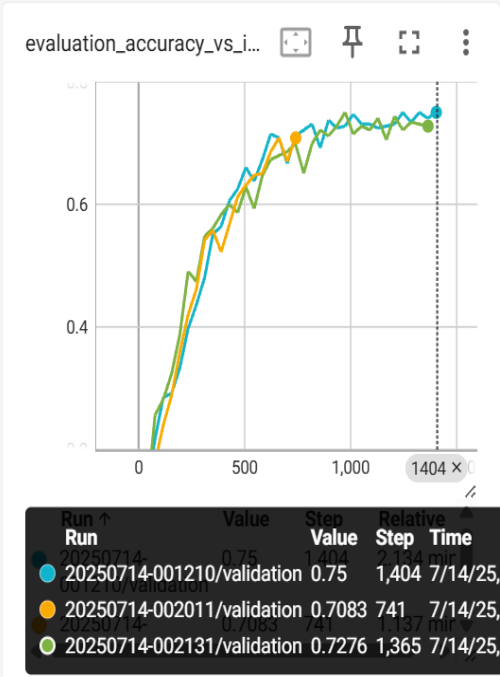
epoch_loss



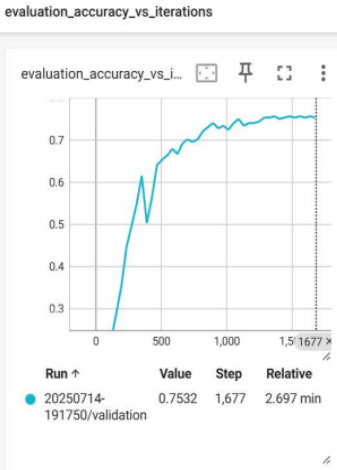

Run	Value	Step	Time	Relative
20250714-001210/train	0.4033	35	7/14/25, 12:14 AM	2.133 min
20250714-001210/validation	1.204	35	7/14/25, 12:14 AM	2.133 min
20250714-002011/train	0.785	18	7/14/25, 12:21 AM	1.137 min
20250714-002011/validation	1.011	18	7/14/25, 12:21 AM	1.137 min
20250714-002131/train	0.3433	34	7/14/25, 12:23 AM	2.131 min
20250714-002131/validation	1.123	34	7/14/25, 12:23 AM	2.131 min

evaluation accuracy vs iterations


evaluation_accuracy_vs_iterations



- 20250714-191750
- 20250714-191750/train
- 20250714-191750/validation



confusion_matrix




- HA2
- aug_search/cnn_best_aug
 - trial_00
 - trial_01
 - trial_02
 - trial_03
 - trial_04
 - trial_05
 - trial_06
 - trial_07
 - trial_08
 - trial_09
 - trial_10
 - trial_11
 - trial_12
 - trial_13
 - trial_14
- oracle.json
- tuner0.json
- eval_artifacts
- Industrial100_Auszug
- ipy_nb_bilder
- tb_final/20250714-191750
 - train
 - validation
- events.out.tfevents.1752513...

Data

- Abstandshalter
- Auslassventil
- Blechlineal
- Filterkartusche
- Gewindestange
- Hohlschraube
- Hutmutter
- Hydraulikstutzen
- Nutenstein
- Schraubenfeder

HA2	tuner_results_sorted.csv	data
1	trial_id,f0,f1,f2,kernel_size,dense_units,dense_dp,weight_decay,optimizer,lr,batch_size,use_plateau,status,best_epoch,train_acc,val_acc	
2	5,128,256,256,7,128,0.3,4.72631111933952e-05,adam,0.0001,16,True,COMPLETED,16,0.96073716878891,0.6570512652397156	
3	38,32,64,256,3,128,0.4,4.578772662888393e-06,adam,0.0005,16,True,COMPLETED,16,0.9631410241127014,0.6538461446762085	
4	24,32,64,256,5,64,0.4,5.844476066973506e-06,adam,0.0005,32,True,COMPLETED,17,0.9054487347602844,0.6538461446762085	
5	17,32,64,128,3,128,0.4,2.743043478381547e-05,adam,0.001,16,False,COMPLETED,30,0.9743589758872986,0.6314102411270142	
6	20,128,64,256,7,128,0.2,9.403614999484184e-06,rmsprop,0.0001,16,False,COMPLETED,16,0.94551283121109,0.6282051205635071	
7	13,128,64,256,3,128,0.3,1.1844905074722785e-06,rmsprop,0.0005,16,True,COMPLETED,18,0.9591346383094788,0.6185897588729858	
8	37,32,64,256,3,64,0.4,1.0322333113531436e-05,adam,0.0005,16,False,COMPLETED,13,0.8790063858032227,0.6153846383094788	
9	39,64,64,256,3,128,0.3,2.9159709503615547e-06,adam,0.0005,32,True,COMPLETED,18,0.9759615659713744,0.6121794581413269	
10	36,64,64,256,5,64,0.4,6.786446896999284e-06,adam,0.0005,32,True,COMPLETED,18,0.8846153616905212,0.5993589758872986	
11	31,128,64,256,3,64,0.4,1.7500478839025897e-05,adam,0.0005,16,True,COMPLETED,20,0.8060897588729858,0.5993589758872986	
12	4,32,64,128,3,128,0.2,2.685272714053744e-05,rmsprop,0.0001,64,False,COMPLETED,25,0.8717948794364929,0.5961538553237915	
13	6,128,64,256,3,256,0.3,2.223834398190109e-06,adam,0.001,32,False,COMPLETED,14,0.99198716878891,0.5961538553237915	
14	7,128,256,256,7,128,0.2,5.105934690849343e-05,sgd,0.001,64,True,COMPLETED,34,0.8862179517745972,0.5961538553237915	
15	18,64,256,256,7,128,0.2,4.729436202637012e-05,rmsprop,0.001,16,True,COMPLETED,26,0.9575320482254028,0.5929487347602844	
16	22,64,128,128,5,256,0.4,5.582095371314753e-06,sgd,0.001,64,False,COMPLETED,34,0.8645833134651184,0.5897436141967773	
17	0,128,64,256,7,64,0.2,3.507215228728303e-05,rmsprop,0.0001,16,False,COMPLETED,20,0.9278846383094788,0.5897436141967773	
18	2,32,128,256,3,64,0.4,8.128524746629793e-05,rmsprop,0.0005,16,False,COMPLETED,24,0.967147409915924,0.557692289352417	
	286746057803035e-06,sgd,0.0005,32,True,COMPLETED,37,0.7323718070983887,0.5865384340286255	
	5114980504743e-06,rmsprop,0.0005,64,True,COMPLETED,15,0.9575320482254028,0.5833333134651184	
	353425444136103e-05,sgd,0.001,32,True,COMPLETED,19,0.9823718070983888,0.5641025900840759	
	2107327077037935e-05,rmsprop,0.001,16,False,COMPLETED,17,0.9054487347602844,0.5641025900840759	
	620301758122103e-05,rmsprop,0.001,32,False,COMPLETED,26,0.967147409915924,0.557692289352417	
	711862216981024e-05,sgd,0.001,32,True,COMPLETED,38,0.75,0.5128205418586731	
	8375332658590165e-06,sgd,0.001,32,False,COMPLETED,37,0.6931089758872986,0.5032051205635071	
	13758360859015e-06,rmsprop,0.0001,64,True,COMPLETED,35,0.5496794581413269,0.4294871687889099	
	1.7887307871975e-06,adam,0.001,32,False,COMPLETED,20,0.8998397588729858,0.4294871687889099	
	83401779428952e-05,sgd,0.0005,16,False,COMPLETED,39,0.4206730723381042,0.4198718070983886	
	499991835674088e-05,sgd,0.0005,16,True,COMPLETED,16,0.2612179517745971,0.3076923191547394	
	7169648733932947e-06,sgd,0.0005,16,False,COMPLETED,15,0.1866987198591232,0.25	
	.935858504752804e-06,sgd,0.0001,32,False,COMPLETED,10,0.1794871836900711,0.2211538404226303	
	35439859879942e-05,sgd,0.0005,16,True,COMPLETED,20,0.1706730723381042,0.0961538478732109	
	.5289905806123778e-05,sgd,0.0001,32,False,COMPLETED,14,0.196314096458085e,0.20833333283662796	
	806871387360406e-05,sgd,0.0005,16,True,COMPLETED,16,0.1530448645353317,0.1923076957464218	
	2.24669700353512e-05,adam,0.001,64,False,COMPLETED,0,0.0929487198591232,0.0961538478732109	
	.715817770883184e-05,rmsprop,0.001,16,False,COMPLETED,0,0.1105769202113151,0.0929487198591232	
	4612244291025357e-05,adam,0.001,64,True,COMPLETED,2,0.0945512801408767,0.0929487198591232	
	1.6899045670308515e-06,adam,0.001,32,False,COMPLETED,0,0.1049679517745971,0.0865384638309478	
	700829010008494e-05,adam,0.0005,64,False,FAILED,0,,	
	2.1816829689460077e-06,rmsprop,0.0001,16,True,FAILED,0,,	
	2.834241861553627e-05,adam,0.0001,64,False,FAILED,0,,	



Classification Report – Final Model

	precision	recall	f1-score	support
Abstandshalter	0.963	0.963	0.963	27
Auslassventil	0.923	0.800	0.857	30
Blechlineal	0.727	0.593	0.653	27
Filterkartusche	0.964	1.000	0.982	27
Gewindestange	0.800	0.800	0.800	35
Hohlschraube	0.667	0.643	0.655	28
Hutmutter	0.759	0.667	0.710	33
Hydraulikstutzen	0.667	0.690	0.678	29
Nutenstein	0.588	0.541	0.563	37
Schraubenfeder	0.630	0.872	0.731	39
accuracy			0.753	312
macro avg	0.769	0.757	0.759	312
weighted avg	0.759	0.753	0.752	312

These things are written above saperately in the code, here is a summary of the process.

Bewertung der Optimierungsschritte (DE)

Die Optimierung erfolgte schrittweise, beginnend mit einem Basismodell und anschließend durch eine umfangreiche Hyperparameter-Suche:

1. **Basismodell:** Das ursprüngliche Modell (3 × [Conv+Pool] + Flatten + Dense) erzielte nach 40 Epochen eine Validierungsgenauigkeit (**val_acc**) von etwa **60,3 %** bei einem Verlust (**val_loss**) von **1,57**. Dieses Ergebnis dient als Referenzpunkt für die weitere Optimierung.

2. **Architektur- und Trainingsoptimierung:** Durch Bayesian-Optimierung der Architektur- und Trainingsparameter (Anzahl der Filter, Kernelgröße, Dropout-Rate, Gewicht-Regularisierung, Lernrate, Batch-Größe und Lernraten-Scheduler) wurde ein deutlicher Fortschritt erzielt. Die besten gefundenen Hyperparameter ($f_0=128$, $f_1=256$, $f_2=256$, Kernelgröße=7, Dense=128, Dropout=0,3, Gewicht-Regularisierung= $4,73e-5$, Adam-Optimizer mit LR= $1e-4$ und Batchgröße=16) führten zu einer signifikanten Steigerung der Validierungsgenauigkeit auf etwa **65,7 %** bei einem Verlust von **1,90**. Die größere Kernelgröße (7) und eine stärkere Modellkapazität trugen dabei entscheidend zur Verbesserung bei.
3. **Optimierung durch Datenaugmentation:** Zusätzlich wurde mittels gezielter Bayesian-Optimierung eine systematische Untersuchung der Datenaugmentation durchgeführt (Rotation, Zoom, horizontale Spiegelung). Die optimalen Einstellungen (Rotation: 10 %, Zoom: 0 %, Flip: False) führten zu einer weiteren signifikanten Steigerung der Validierungsgenauigkeit auf rund **79,2 %** und zu einem reduzierten Verlust von **1,05**. Besonders die moderate Rotation erwies sich hierbei als wirksam, um die Robustheit des Modells zu verbessern.
4. **Trainingstaktiken:** Der Einsatz von Early-Stopping mit Geduld von 10 Epochen sowie eines Plateau-LR-Schedulers (Reduzierung der Lernrate bei Stagnation) verbesserten zusätzlich die Stabilität und Konvergenz des Trainings deutlich, reduzierten das Overfitting-Risiko und gewährleisteten einen optimalen Trainingsabschluss.

Fazit (Gesamtbewertung)

Insgesamt brachte die systematische und iterative Optimierung – bestehend aus Architektur- und Hyperparameter-Suche sowie der gezielten Datenaugmentation – eine substantielle Steigerung der Modelleleistung. Die Validierungsgenauigkeit erhöhte sich insgesamt um rund **18,9 Prozentpunkte** gegenüber dem ursprünglichen Basismodell. Die sorgfältige Kombination aus Modellstruktur, Trainingsmethoden und Datenaugmentation erwies sich somit als äußerst effektiv für die Optimierung.

below is english version (for myself).

Assessment of the Optimization Steps (EN)

The optimization proceeded step-by-step, starting from a baseline model and followed by extensive hyperparameter tuning:

1. **Baseline Model:** The initial model ($3 \times [\text{Conv}+\text{Pool}] + \text{Flatten} + \text{Dense}$) achieved a validation accuracy (**val_acc**) of approximately **60.3%** with a loss (**val_loss**) of **1.57** after 40 epochs. This served as the baseline for subsequent optimizations.
2. **Architecture and Training Optimization:** Bayesian optimization of architectural and training parameters—including filter sizes, kernel size, dropout rates, weight decay, learning rate, batch size, and the learning rate scheduler—resulted in significant improvements. The optimal hyperparameters ($f_0=128$, $f_1=256$, $f_2=256$, kernel size=7, dense units=128, dropout=0.3, weight decay= $4.73e-5$, Adam optimizer with LR= $1e-4$, and batch size=16) notably improved validation accuracy to about **65.7%**, although with a slightly higher loss (**1.90**). Larger kernel sizes (7) and increased model capacity played a critical role in this improvement.

3. **Optimization via Data Augmentation:** Additionally, a systematic Bayesian optimization of data augmentation parameters (rotation, zoom, horizontal flip) was conducted. The optimal settings (rotation: 10%, zoom: 0%, flip: False) further significantly increased validation accuracy to around **79.2%** while simultaneously reducing loss to **1.05**. Moderate rotation proved particularly beneficial in enhancing the robustness of the model.
4. **Training Techniques:** Incorporating Early-Stopping (patience=10) and a Reduce-LR-on-Plateau scheduler significantly improved training stability and convergence, reduced the risk of overfitting, and ensured optimal model training completion.

Conclusion (Overall Assessment)

Overall, the systematic and iterative optimization—consisting of architectural and hyperparameter search combined with targeted data augmentation—yielded substantial improvements in model performance. The validation accuracy increased by approximately **18.9 percentage points** compared to the original baseline model. The careful combination of model architecture, training strategies, and data augmentation proved highly effective in optimizing performance.

I Think

The final validation accuracy of around 79% is considered acceptable in this case.

This accuracy is because the model is designed to classify 10 different industrial object categories, including visually similar and easily confusable items such as "Blechlineal" and "Hutmutter." According to the classification report, most classes show precision and recall values is higher than 0.75, with only a few categories like "Blechlineal, Nutenstein" performing lower.

For applications that require higher performance or stricter accuracy, further improvements are possible.

These could include using larger model architectures such as ResNet, applying more advanced data augmentation techniques like MixUp, or converting images to grayscale to reduce irrelevant color variations.